



This document details the procedures to set up the docker containers from NGC.

## Prerequisites

- If using a new server with NVIDIA GPU without having NVIDIA drivers, CUDA toolkit, Docker CE software, Nvidia-container runtime and NVIDIA Docker stack installed, the below steps need to be completed, assuming you are using **Ubuntu OS**.
- If you do not have an Nvidia GPU, at least have docker 19.03.

**(DO NOT RUN THE BELOW STEPS IF YOU HAVE THE ABOVE STACK ALREADY INSTALLED)**

- **Install NVIDIA Drivers**

Download and install the relevant drivers (Tesla V100 / Tesla T4) from

<https://www.nvidia.com/Download/index.aspx>

**OR (use the below steps - Preferred)**

- `sudo apt-get purge nvidia*`
- `sudo add-apt-repository ppa:graphics-drivers/ppa`
- `sudo apt update`
- `sudo apt install nvidia-driver-450`

**Note:** Any driver starting nvidia-driver-410 or later will work.

- **Install CUDA drivers**

Step 1: Check the currently installed release

- `apt list --installed cuda-toolkit-*`

Step 2: Update the local database with the latest information from the Ubuntu repository.

- `sudo apt update`

Step 3: Show all available CUDA Toolkit releases.

- `apt list cuda-toolkit-*`

Step 4: Install or upgrade the CUDA Toolkit

- `apt install cuda-toolkit-11-2`

After performing the above steps (either of the above) **reboot** the instance and check the status of the GPU by using the below:

➤ *nvidia-smi*

NVIDIA-SMI 451.67			Driver Version: 451.67			CUDA Version: 11.0		
GPU	Name		TCC/WDDM	Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.	
0	Quadro M1200		WDDM	00000000:01:00.0	Off			N/A
N/A	42C	P8	N/A / N/A		40MiB / 4096MiB	0%	Default	
Processes:								
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage		
	ID	ID						
No running processes found								

- **Install docker and NVIDIA docker2:**

Check if already installed: ➤ ***docker version***

If Not, Install Docker CE using the below steps:

<https://docs.docker.com/engine/install/ubuntu/#install-using-the-repository>

- *sudo apt-get update*
- *sudo apt-get install apt-transport-https ca-certificates curl gnupg-agent software-properties-common*
- *curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -*
- *sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu \$(lsb\_release -cs) stable"*
- *sudo apt-get update*
- *sudo apt-get install docker-ce docker-ce-cli containerd.io*

ONLY FOR NVIDIA GPU USER:-

**Install nvidia-docker2 using the below steps:**

<https://docs.nvidia.com/datacenter/cloud-native/container-toolkit/install-guide.html#installing-on-ubuntu-and-debian>

- *curl https://get.docker.com | sh && sudo systemctl start docker && sudo systemctl enable docker*
- *distribution=\$(cat /etc/os-release; echo \$ID\$VERSION\_ID) && curl -s -L https://nvidia.github.io/nvidia-docker/gpgkey | sudo apt-key add - && curl -s -L https://nvidia.github.io/nvidia-docker/\$distribution/nvidia-docker.list | sudo tee /etc/apt/sources.list.d/nvidia-docker.list*
- *curl -s -L https://nvidia.github.io/nvidia-container-*

```
runtime/experimental/$distribution/nvidia-container-runtime.list | sudo tee
/etc/apt/sources.list.d/nvidia-container-runtime.list
```

- `sudo apt-get update`
- `sudo apt-get install -y nvidia-docker2`
- `sudo systemctl restart docker`

Verify the installation of docker and nvidia-docker with following commands:

- `sudo docker run --rm --gpus all nvidia/cuda:11.0-base nvidia-smi`
- `sudo docker run --rm hello-world`

## Step wise commands to run jupyter-lab

`docker pull frodo/alpine-python-machinelearning`

1. Show the running container

```
$ sudo docker ps
```

2. Show the docker images

```
$ sudo docker images
```

3. Running the container image

- a. If docker version  $\geq 19.02$

General command:

```
$ docker run --gpus all --rm -v /path/to/data:/workspace/data -it -p 5000:8888
nvcr.io/nvidia/tensorflow:20.03-tf1-py3
```

Example:

→ Highlighted in red, needed to be change

```
$ docker run -it --gpus device=0 --rm -v $PWD/data:/workspace/data -p 5000:8888
nvcr.io/nvidia/tensorflow:21.02-tf2-py3
```

`--gpus device = 0`

Based on the gpu device you have to use, you can give device as 0, 1, 2, 3(n-1 = device)

`-v $PWD/data:/workspace/data`

Your present working directory, then add the directory you want to map

e.g. `/home/dgx/nvidia` → so, docker workspace will be stored in the nvidia directory.

```
-p 1000:8888
```

Here, 8888 is the default port and 1000 is the port you want to forward your default port to.

#### 4. Running Jupyter-Notebook

- Execute a JupyterLab-

Install Jupyter lab

```
$ pip install jupyterlab
```

Run jupyter lab in the docker workspace

```
$ jupyter-lab --allow-root --ip=0.0.0.0
```

Note: Copy the token and type IP and port assigned (see docker run) on your browser

a. If on your own system-

localhost:5000

Click Enter and paste token

b. If on a different server IP connected with VPN:-

a.b.c.d:5000

Click Enter and paste token

\*\*\*

## Additional Commands

If you have container running in the background, then you can follow these steps:

First get the container detail

```
$ docker ps -a
```

Method 1: Stop and remove the running container

```
$ docker stop container_id
```

```
$ docker rm container_id
```

Method 2: Start and attach the running container

```
$ docker start container_id
```

```
$ docker attach container_id
```